

PRODUCT SYSTEM OF RECORD

The Vocari Book

Product vision, current-state architecture, approved direction, and responsible evolution

NORTH STAR

Help people recognize recurring patterns in their lived experience, understand what those patterns may mean for work, and choose an evidence-informed next step without being reduced to a personality label.

Version 1.0 | 6 August 2026 | Decision guide for product, design, engineering, content, and AI collaborators

Prepared from the live Lovable project, repository snapshot, final course research, Progressive Discovery Engine specifications, and Career Content & Data Specification v2.1.

How to use this book

This is a decision boundary, not a feature wish list.

The documents behind Vocari contain strong ideas, but they were written at different moments and for different purposes: academic research, prototype handoff, content governance, implementation planning, and live product changes. This book reconciles them into one operating model. It prevents a future designer, engineer, Lovable prompt, or AI agent from treating a prototype claim as a shipped capability.

Label	Meaning	Decision rule
CURRENT REALITY	Verified in the August 2026 code snapshot or live project metadata.	May be described as available, while preserving limitations stated here.
APPROVED DIRECTION	Supported by research or an approved specification, but not fully implemented.	May guide near-term work. Must be broken into testable increments before being marketed.
LONG-TERM VISION	Strategically consistent idea with unresolved product, evidence, business, or architecture questions.	May inform architecture seams and research. Must not be presented as committed.
OUT OF SCOPE / HOLD	Attractive idea that would distract from the validated core or creates disproportionate risk.	Do not build without a new decision record and evidence.

Decision precedence

1. Safety, privacy, user agency, and legal/ethical constraints.
2. This book and later approved decision records.
3. Career Content & Data Specification v2.1 for career facts and matching semantics.
4. Progressive Discovery Engine specifications for profile learning behavior.
5. Live code and validated data for statements about current capability.
6. Older handoffs, course drafts, screenshots, and chat messages as historical evidence, not automatic requirements.

NON-NEGOTIABLE

If copy says Vocari “learns,” “remembers,” “improves,” or “understands,” the product must identify the stored evidence, the transformation applied, the user-visible benefit, and the user’s control over the change.

Executive product definition

Vocari is a reflective career discovery platform for people whose experiences, strengths, interests, and values do not collapse neatly into one label. Its product is self-understanding. Career recommendations are an applied output of that understanding, not the entire value proposition.

Element	Canonical statement
Mission	Help people recognize evidence-based patterns in their experiences and translate those patterns into confident, realistic career exploration.
Vision	A trusted, evolving career-development companion that helps users understand themselves before choosing, changing, or developing a career.
Primary problem	Career uncertainty is both practical and emotional. Generic assessments create answers without sufficient explanation, context, or user agency.
Primary audience	Multipotentialites, career changers, graduate and recent graduates, and capable professionals with nonlinear paths or decision paralysis.
Not a job board	Vocari explains careers and fit. It does not optimize for vacancies, applications, or fastest placement.
Not an authority	Vocari offers an interpretation of evidence. The user remains the expert on their life and may clarify or challenge the interpretation.
Desired outcome	The user can name patterns, explain why a career is worth exploring, identify a realistic tradeoff, and choose a next step with greater confidence.

The product promise and value exchange

Every meaningful contribution should return immediate value. A user who adds a reflection should not be asked to trust an invisible future benefit. Vocari should show what was learned, where the evidence came from, what changed, what did not change, and which next question would most improve understanding.

User contribution	Immediate return	What must not happen
Foundation answer	A named early pattern plus evidence source.	A definitive identity label or unexplained percentage.
Reflection	A new or strengthened signal and a transparent effect.	A saved note that never affects any insight.
Career exploration	A clearer comparison or optional preference signal.	Silent behavioral surveillance or drastic ranking change.
Clarification or challenge	A proposed profile refinement requiring confirmation.	Silent profile rewriting by AI.
Saved career	A preserved exploration set and better next-step prompts.	Treating a save as proof of enduring fit.

1. Evidence and research foundation

What the source material supports, and what it does not prove.

The final course work converged on three needs: personalized guidance, emotional support during career exploration, and transparent explanations that improve trust. Interviews, surveys, personas, perception testing, experience testing, behavioral observation, journey mapping, and limited physiological research pointed in the same broad direction. This convergence is strong enough to guide product priorities, but the sample sizes and prototype conditions do not prove clinical efficacy, algorithmic validity, retention, or market demand.

Evidence-supported conclusion	Strength	Boundary
Users value guidance connected to their own experiences.	Strong qualitative convergence.	Does not validate any specific scoring algorithm.
Explanation matters when results surprise or conflict with expectations.	Observed in comments and hesitation behavior.	More explanation can still be wrong if the career data or inference is weak.
Career exploration includes anxiety, comparison, and fear of a wrong choice.	Repeated across methods and personas.	Vocari is not therapy and should not make mental-health claims.
Users can be cognitively engaged while feeling only moderately supported.	Self-report and physiological findings diverged in a useful way.	MorphCast data came from only a subset and should not drive personalization.
A progressive model fits the stated need better than a one-time result.	Strong conceptual fit and user interest.	Longitudinal benefit has not yet been demonstrated.

Research constraints to preserve

- Eight participants completed the prototype in the affective-computing work; usable facial analytics were available for three. This is exploratory evidence, not a universal emotional model.
- Testing used a prototype with a limited career dataset. Result mismatch may reflect catalog coverage as much as matching quality.
- High reported ease of use coexisted with hesitation, rereading, and repeated clicks. Behavioral friction should not be dismissed because a completion score is positive.
- Participants asked for more encouragement, but emotional support should be specific and evidence-based, not empty praise or dependency-forming gamification.
- Future validation needs longitudinal use, broader occupations and education paths, larger samples, accessibility testing, and measurement of understanding rather than only satisfaction.

2. Experience principles

Principle	Experience requirement	Failure signal
Recognition before prescription	Surface patterns with evidence before presenting careers.	The first meaningful outcome is a ranked job list.
Interpretation, not classification	Use tentative language and allow correction.	“This is who you are” or fixed personality types.
Explainability at the point of doubt	Connect recommendations to user signals and career facts.	Generic “why it fits” text reused across users.

Principle	Experience requirement	Failure signal
Stable growth	Small new evidence deepens explanation before reordering careers.	One answer radically changes the ranking.
Agency by design	Ask before applying profile refinements; allow editing and deletion.	A model silently overwrites a user profile.
Tradeoffs are part of fit	Every strong recommendation names a realistic friction.	Only flattering evidence is shown.
Progress without grading	Show what is known and what remains unclear.	A low “confidence score” feels like failure.
Accessible exploration	Career browsing is available without an account.	Personalization becomes a gate to basic career facts.

TONE STANDARD

Warm, calm, specific, and realistic. Vocari can validate effort and uncertainty, but it should not tell a user that a career is “meant for” them, guarantee fulfillment, or imply that continued use is necessary to know themselves.

3. Current reality: product surface

Verified from the repository snapshot and live project as of 6 August 2026.

Surface	Current capability	Important limitation
Homepage	Recognition-led positioning, discovery CTA, sample dashboard, progressive-profile story.	Marketing language can outrun the actual learning loop.
Onboarding	Five sections: context, interests, paired preferences, work-style sliders, and reflections; before/after confidence rating.	Foundation flow remains a bounded assessment even when framed as ongoing.
Results	Rule-derived patterns, evidence snippets, tiered recommendations, and explanation surfaces.	Free text is not semantically analyzed; patterns are deterministic.
Explore	Public career library with search, filters, salary/time-to-entry controls, and career cards.	Client-side catalog; factual quality varies by migrated legacy record.
Career detail	Day-to-day, fit, environment, skills, progression, similar careers, and personalized compatibility when signals exist.	Some content is adapter-generated or low-confidence; “AI impact” is content, not an AI feature.
Dashboard	Completion state, pattern cards, recommendations, snapshot, and a reflection journal.	Dashboard reflection entries remain local and do not currently feed profile inference.
Authentication	Email/password flows and password reset through Supabase.	No complete user data-management center is visible.
Persistence	Profile is localStorage-first and syncs to a per-user JSONB profile for signed-in users.	Cloud update assumes an existing profile row; conflict resolution is cloud-wins, not version-aware.
Research instrumentation	Consent-aware MorphCast loader, server-side emotion capture, and admin research view.	Research tooling is not a consumer personalization engine and should be separated or retired after research.
Demo	A sample persona dashboard shows the intended emotional payoff.	Demo content must never be confused with a computed user result.

4. Current reality: architecture and data flow

The implementation is a TanStack Start / React 19 TypeScript application with file-based routes, Tailwind, shadcn/Radix components, Supabase authentication and persistence, and static TypeScript career content. The architecture can support a credible prototype and incremental production hardening. It does not yet support an always-learning AI companion without new event, evidence, versioning, and conversation layers.

Layer	Implemented now	Architectural implication
Presentation	Route components plus shared UI and recommendation components.	Good for iteration; domain logic should continue moving out of route files.
User profile	One nested Profile object saved locally and as JSONB.	Fast, but coarse-grained updates, provenance, migrations, and auditability are difficult.
Signal derivation	Deterministic mappings from selected answers and fields to normalized signals.	Explainable and testable, but limited vocabulary and no semantic free-text understanding.
Pattern derivation	Rule-based functions with evidence strings.	Current personalization must be described as rules, not AI interpretation.

Layer	Implemented now	Architectural implication
Matching	Six weighted dimensions and career Profile DNA comparisons with confidence separated from score.	Semantically aligned with v2.1; needs calibration and validation before strong accuracy claims.
Career catalog	24 slugs in one legacy file, including a duplicate senior UX Researcher pathway, plus v2.1 schema and compatibility adapter.	Migration is underway, not complete. Catalog goal is about 75 balanced careers.
Backend	Profiles, roles, emotion sessions/events, RLS policies, service-side research writes.	No discovery-event ledger, conversation tables, saved-career store, or roadmap state yet.
AI services	No production conversational inference path in the supplied snapshot.	v1.1 AI Gateway plan is approved direction, not current reality.

Current data flow

Onboarding input → local Profile object → optional JSONB sync → deterministic user-signal derivation → rule-based patterns → compatibility scoring against adapted career records → Results, Dashboard, and Career Detail presentation.

ARCHITECTURE TRUTH

The current app can reasonably support deterministic progressive discovery first. It cannot honestly promise that every reflection, page view, or conversation improves recommendations until those events are stored with provenance and consumed by a versioned inference process.

5. Career content system: approved contract

Career Content & Data Specification v2.1 is the canonical career-content contract. It deliberately allows TypeScript or validated JSON for the academic prototype and does not require a CMS. This is the correct constraint: structured, governed content is more urgent than platform migration.

Contract area	Approved rule
Separation of concerns	Career facts are global; user match scores and personalized explanations are calculated at runtime.
Publication states	Draft may be partial. Published requires complete fields, validation, human review, verification dates, and valid relationships.
Profile DNA	Twenty independent 0–100 career-characteristic intensities; high is not inherently better and values do not total 100.
Recommendation weights	Six career-specific weights total 100: interests, work style, values, motivators, strengths, and environment.
Confidence	Career evidence confidence describes record quality. Recommendation confidence describes evidence for one user. Neither is the match score.
Sources	Compensation, outlook, licensing, and education claims require claim-level source paths and access dates.
Taxonomy	Industries, enums, and tags are controlled. UI labels may vary; stored values may not drift.
Relationships	Immutable career IDs connect similar paths, transitions, progression, and specializations.
Search	OR within a category and AND across categories; factual constraints remain distinct from exploratory preferences.
Scale	Client-side filtering remains reasonable near 75 careers; a database/CMS is an operational decision, not a prerequisite.

Migration reality

- The repository includes the v2.1 Zod schema, taxonomies, Profile DNA dimensions, adapter, filters, presentation mapping, and compatibility logic.
- The legacy catalog contains 24 slugs. The approved migration calls for merging ux-researcher-2 into the canonical UX Researcher record and preserving the retired slug as an alias.
- Adapter defaults keep the app functional but produce low-confidence inferred values. They are migration scaffolding, not publication-quality content.
- The target of approximately 75 careers should emphasize balanced industry, education, and non-degree pathway coverage, not numerical expansion alone.

6. Recommendation system: truth and target

Aspect	Current reality	Approved direction
Inputs	Selections, slider values, paired choices, demographics, and selected reflection values; some free-text is used for gating/evidence presence.	Append-only, source-attributed discovery signals from foundation, sessions, reflections, confirmed conversations, and deliberate exploration actions.

Aspect	Current reality	Approved direction
Scoring	Deterministic dimension compatibility using adapted career Profile DNA and weights.	Versioned matching contract with calibration tests, missing-signal handling, stability rules, and consistent consumers.
Explanation	Matched dimensions, evidence snippets, compatibility notes, and career-specific content.	Trace each explanation to user evidence and governed career facts; disclose weak or approximate evidence.
Confidence	Separated in code from match score.	Explain as evidence sufficiency, never as a grade; show what would materially improve understanding.
Stability	Some tier gating and confidence adjustment exist.	Persist result snapshots and compare meaningful changes so careers do not silently reshuffle.
User control	Profile sections can be edited.	Let users inspect, correct, confirm, or decline profile refinements and see downstream effects.

ACCURACY LANGUAGE

Until a broader catalog, content review, and calibration study exist, use “recommendations for exploration” and “evidence-informed fit,” not “accurate prediction,” “best career,” or “scientifically validated match.”

7. Approved direction: Progressive Discovery Engine v1

The approved near-term product direction is to make progressive learning real through small, optional Discovery Sessions. This is more credible than jumping directly to an unconstrained AI coach because it can reuse deterministic signals, preserve explainability, and validate whether ongoing contribution actually creates value.

Capability	Minimum viable behavior	Required dependency
Discovery event ledger	Store each meaningful input as an append-only event with source, timestamp, profile version, and evidence link.	New backend schema and migration strategy.
Coverage model	Identify which signal domains are strong, thin, contradictory, or stale.	Versioned signal taxonomy and evidence counts.
Session selector	Choose 2–4 prompts that fill high-value evidence gaps.	Question bank mapped to signal domains; no AI required initially.
Immediate payback	After a session, state what was learned and what explanation deepened.	Before/after inference comparison and user-facing change summary.
Stable recommendations	Small evidence usually changes confidence or explanation before rank.	Persisted result snapshots and threshold rules.
Contradiction handling	Describe changed or context-dependent priorities without calling the user inconsistent.	Evidence chronology and contextual qualifiers.
Understanding progress	Show coverage and clarity in words or bands, not a grade-like percent.	Accessible status vocabulary and tested copy.
Progressive prompt memory	Reference earlier confirmed evidence when asking the next question.	Evidence retrieval with privacy controls.

Recommended delivery sequence

7. Define the event and evidence contracts before changing copy.
8. Persist foundation answers, dashboard reflections, saves, and Discovery Session responses in one source-attributed ledger.
9. Implement deterministic gap selection and a small curated question bank.
10. Show evidence counts and an explicit change summary after each session.
11. Add result snapshots and stability thresholds.
12. Run longitudinal usability testing before expanding inputs or adding generative conversation.

8. Approved direction: Collaborative Profile Refinement v1.1

The recent Lovable discussion approved a narrow first conversation scope: Profile DNA traits first, Lovable AI Gateway first, and backend persistence with authentication required. Lovable produced a phased implementation plan, but no implementation is present in the supplied repository snapshot. Therefore this is approved direction, not current capability.

Boundary	Approved choice	Why it matters
Entry point	A specific Profile DNA trait, never a blank chatbot.	Keeps the conversation purposeful and limits hallucinated scope.
Modes	Understand, Clarify, Challenge.	Separates explanation from user-supplied correction.

Boundary	Approved choice	Why it matters
Interaction	Guided chips plus free text, generally 3–5 turns.	Collects structured evidence while leaving room for nuance.
Authority	Vocari proposes; the user confirms or declines.	Prevents silent AI ownership of identity.
Persistence	Authenticated, user-owned conversation and turn records with RLS.	Conversations contain sensitive autobiographical data.
Output	Plain-language learnings plus bounded proposed adjustments.	Makes the effect inspectable before application.
Integration	Confirmed refinements become source-attributed discovery signals.	Preserves the existing matching pipeline and provenance.
Payback	Show what changed in understanding and any affected explanation.	Maintains the contribution-to-value exchange.

Architecture guardrails

- Keep the AI key and system instructions server-side. Validate structured outputs against a strict schema.
- Do not let a model write directly to the canonical profile or career records. Convert confirmed proposals into bounded domain events.
- Store prompt/model/algorithm versions and the evidence context supplied to the model.
- Redact or minimize unnecessary sensitive content; define retention, export, and deletion behavior before launch.
- Treat conversation summaries as user-confirmed interpretations, not objective facts.
- Start with traits only. Recommendations, roadmaps, values, reflections, and open coaching remain later experiments.

9. Contradictions and unresolved concepts

These must be resolved explicitly rather than hidden in copy.

Contradiction	Why it matters	Resolution
“Not a personality assessment” vs. a section called personality and trait sliders.	The interaction can undermine the positioning.	Rename the internal/user-facing concept to work tendencies or preference signals; document exactly how it differs from typology.
“Every reflection sharpens the picture” vs. dashboard reflections stored separately in localStorage.	The promise is currently false for journal entries.	Either integrate reflections into the signal ledger or soften the promise until integration ships.
“Profile DNA evolves” vs. rule-derived patterns without durable evidence history.	Users cannot inspect change or provenance.	Create versioned evidence and profile snapshots before marketing evolution.
Limited catalog vs. confident career percentages.	A good match may be absent, and ranking among 23 unique pathways is not comprehensive.	Disclose coverage, avoid “best,” and expand balanced careers after governance migration.
Profile DNA describes careers, while the UI also uses Profile DNA for the user.	The same term can mean two different vectors.	Adopt explicit names: Career DNA and User Profile DNA, or reserve Profile DNA for the user and name the career vector Career Characteristics.
Career Compass vs. Vocari naming persists across storage keys, docs, and metadata.	Creates implementation and analytics ambiguity.	Treat Career Compass as historical working title; migrate keys carefully and keep compatibility aliases.
10 vs. 15 minute duration statements.	Damages trust at the first commitment.	Measure actual completion time and publish a range or stage-specific estimate.
Confidence as emotional self-rating, evidence confidence, and recommendation confidence.	One word represents three constructs.	Name them separately: career confidence, evidence strength, and recommendation evidence level.
Research emotion capture vs. consumer product.	Users may assume emotion is part of personalization.	Keep research instrumentation isolated, opt-in, time-bounded, and absent from recommendation inputs.
Roadmap language vs. current career progression content.	A progression list is not a personalized roadmap.	Use “career progression” now; reserve “personal roadmap” for persisted goals, constraints, milestones, and revision history.

10. Missing concepts required for a trustworthy product

Missing concept	Minimum decision needed	Blocks
Identity vocabulary	Define user Profile DNA vs. career characteristic vectors and every confidence construct.	Consistent UI, analytics, and AI prompts.
Discovery event contract	Event types, source IDs, timestamps, consent, deletion, and versioning.	Progressive learning and evidence display.
Profile versioning	How confirmed signals produce a version and how a user reviews/reverts changes.	Stable recommendations and conversations.
Result snapshot contract	What is stored for comparison and what counts as a meaningful change.	Transparent recommendation evolution.
Data lifecycle	Retention, export, account deletion, anonymous-to-auth merge, and device conflict rules.	Production readiness and user trust.
Content operations	Who drafts, verifies, reviews, publishes, and refreshes each career.	Credible expansion to 75 careers.

Missing concept	Minimum decision needed	Blocks
Evaluation model	Metrics for understanding, explanation quality, calibration, coverage, and harmful mismatch.	Claims of improvement or accuracy.
Accessibility baseline	Keyboard, screen reader, contrast, cognitive load, plain language, and motion standards.	Inclusive launch.
Safety boundaries	Crisis language, discrimination, disability, immigration/licensing, financial and education-cost claims.	Responsible guidance.
Business model boundary	Whether monetization changes ranking, access, referrals, or institutional visibility.	Trust and B2B/B2C strategy.

11. Long-term vision: capability horizon

Long-term ideas remain valuable when they are framed as hypotheses and sequenced behind foundations. The architecture can grow toward them, but the current repository should not be distorted to simulate capabilities it does not yet possess.

Horizon	Capability	Preconditions	Status
Near	Discovery Sessions and evidence-based profile evolution.	Event ledger, versioning, prompt bank, change summaries.	Approved direction
Near	Trait-scoped collaborative refinement.	Auth, conversation schema, AI gateway, confirmation flow, safety review.	Approved direction
Near	Governed 75-career catalog and unified result contract.	v2.1 migration, human review, source refresh plan, calibration.	Approved direction
Mid	Personal career roadmap with goals, constraints, milestones, and revisions.	User-state model, roadmap ontology, progress logic, clear non-guarantee language.	Long-term vision
Mid	Skill-gap analysis and learning recommendations.	Reliable career skill data, user capability evidence, provider neutrality, cost/access metadata.	Long-term vision
Mid	Career comparison and scenario planning.	Comparable facts, tradeoff model, saved state, accessibility.	Long-term vision
Mid	Resume and portfolio guidance tied to confirmed evidence.	Document privacy, provenance, user review, no fabricated achievements.	Long-term vision
Far	Mentor matching or community.	Trust and safety, identity verification, moderation, safeguarding, marketplace operations.	Hold until strategy
Far	General AI career coach.	Proven contextual conversations, safety evaluation, cost model, escalation and deletion controls.	Hold until evidence
Far	B2B institutional product.	Consent boundaries, user/organization data separation, bias audit, procurement and accessibility readiness.	Open strategic question

Features to resist for now

- Badges, streaks, or pressure-based notifications that turn self-discovery into compliance.
- Job scraping, one-click applications, or vacancy ranking that converts Vocari into a generic job board.
- Emotion recognition as a hidden personalization signal.
- Autonomous profile changes or career decisions made by an AI agent.
- A community feed before moderation and purpose are defined.
- A CMS migration undertaken only because the catalog is growing; validated source files remain sufficient near the stated scale.

12. Product architecture target

Domain service	Responsibility	Reasonable next implementation
Profile service	Current confirmed user state and versions.	Normalize profile fields; preserve local-first only as offline cache, not source of truth for signed-in users.
Discovery ledger	Append-only evidence and interaction events.	Postgres tables with RLS, source taxonomy, soft deletion, and export identifiers.

Domain service	Responsibility	Reasonable next implementation
Understanding engine	Derive normalized signals, coverage, contradictions, and proposed changes.	Keep deterministic core; add model-assisted extraction only behind schemas and confirmation.
Career repository	Validated career facts, taxonomies, relationships, and governance.	Complete v2.1 migration in TypeScript/JSON before evaluating CMS.
Recommendation engine	Versioned match computation, tradeoffs, confidence, stability, and explanations.	One result contract used by Results, Dashboard, Career Detail, and comparison.
Conversation service	Contextual turns and proposed refinements.	Trait-only authenticated service with server-side AI and bounded outputs.
Roadmap service	Goals, constraints, milestones, actions, and revisions.	Do not implement until the profile/recommendation contracts are stable.
Research analytics	Separate consented study instrumentation.	Feature-flag, isolate, document retention, and retire when study ends.

PRAGMATIC ARCHITECTURE

Do not replace deterministic matching with a language model. Use models where language interpretation adds value, then convert outputs into validated, user-confirmed domain objects that the existing explainable engine can consume.

13. Decision roadmap and release gates

Sequence	Outcome	Release gate
0. Vocabulary freeze	One glossary for profile, evidence, score, confidence, career DNA, and recommendation.	Copy and TypeScript types agree; contradictions listed in this book are resolved.
1. Career integrity	23 unique current pathways migrated and reviewed; duplicate resolved; coverage map created.	All published records pass v2.1 validation and factual claims have sources.
2. Unified recommendations	All product surfaces consume one versioned result contract.	Same user/career pair yields the same score, evidence, tradeoff, and confidence everywhere.
3. Evidence foundation	Discovery events, profile versions, result snapshots, export/deletion, and merge rules.	A user can see and remove evidence; changes are attributable and reversible.
4. Progressive sessions	Optional 2–4 prompt sessions fill identified gaps and return value.	Longitudinal test shows improved understanding without ranking volatility.
5. Trait conversations	Understand/Clarify/Challenge for Profile DNA traits.	Security, privacy, structured-output, confirmation, and harmful-response tests pass.
6. Roadmap pilot	Goals and next steps become persistent and revisable.	Milestones are realistic, sourced, and never presented as guaranteed outcomes.
7. Scale decision	Choose CMS, search service, or B2B capabilities only from operating evidence.	Measured editorial/search limits justify added platform complexity.

Definition of done for any new feature

- The feature is labeled Current, Approved, Vision, or Hold in the decision record.
- User value and the contribution-to-return exchange are explicit.
- Data inputs, outputs, provenance, retention, deletion, and consent are defined.
- Failure states, harmful interpretations, and accessibility have acceptance criteria.
- Copy does not make a stronger claim than the implementation or evidence supports.
- At least one metric tests understanding or decision quality, not only clicks.
- AI output, if any, is schema-validated, bounded, reviewable, and never silently authoritative.
- Documentation and canonical types are updated in the same release.

14. Measurement framework

Goal	Primary measure	Guardrail
Clarity	User can accurately explain why a pattern or career appeared.	Do not count explanation length as understanding.
Confidence	Change in self-rated career decision confidence over time.	Do not pressure users toward artificially higher ratings.
Recommendation trust	Agreement that evidence and tradeoffs feel relevant and inspectable.	Track surprise and disagreement, not only approval.
Progressive value	Share of sessions that produce a user-confirmed useful insight.	Avoid optimizing session frequency or streaks.

Goal	Primary measure	Guardrail
Stability	Rate and magnitude of recommendation changes per meaningful new evidence.	Monitor stagnation as well as volatility.
Coverage	Industries, education levels, settings, and pathways represented in published catalog.	Do not treat raw career count as quality.
Content integrity	Percent of factual claims in date, reviewed, and source-supported.	Block publication when critical claims are unsupported.
Accessibility	Task completion and comprehension across assistive and low-literacy contexts.	Include keyboard and screen-reader defects in release gates.
Safety	Rate of harmful, deterministic, stereotyped, or unsupported guidance.	Manual review of high-risk AI and content cases.

15. Guidance for Lovable, engineers, designers, and AI agents

Role	Use this book to...	Do not...
Lovable / coding agent	Check status labels before planning; inspect canonical types; propose phased changes; preserve evidence and user agency.	Implement a vision feature because it appears in an old handoff or homepage promise.
Engineer	Keep domain contracts versioned; write migration and validation tests; maintain deterministic seams around AI.	Let route components or prompts become the only source of business logic.
Designer	Design for explanation, correction, tradeoffs, and uncertainty; test comprehension and emotional support together.	Use confidence percentages, celebratory pressure, or flattering copy as substitutes for evidence.
Content designer	Apply controlled vocabulary, plain language, dated facts, and realistic tradeoffs.	Create generic repeated “why it fits” paragraphs or unsupported salary/outlook claims.
Researcher	Separate prototype usability, algorithm validity, longitudinal impact, and market demand.	Generalize small exploratory physiological samples to all users.
Product owner	Maintain the status map, decision log, and release gates; remove promises when implementation changes.	Allow business experiments to compromise ranking neutrality or profile privacy without a decision record.
AI model	Treat the user as authority; cite available evidence; ask bounded questions; label uncertainty; propose rather than apply.	Diagnose, stereotype, fabricate user history, or guarantee outcomes.

Canonical prompt preamble

FOR FUTURE AGENTS

Before proposing or changing Vocari, classify the request as current maintenance, approved-direction implementation, or long-term exploration. Verify the repository and canonical specification. State which user evidence and career facts the change consumes, how the user can inspect or correct it, and what claim the UI is permitted to make after shipping.

16. Governance and decision log template

Field	Required content
Decision ID	Stable ID, date, owner, and affected release.
Status	Proposed, approved, implemented, validated, superseded, or retired.
Problem	Observed user or system problem, not a desired feature.
Evidence	Research, analytics, defect, content audit, or architectural constraint.
Decision	The chosen direction and explicit non-goals.
Architecture	Affected contracts, tables, routes, privacy, security, and migration.
Experience	User value, copy boundary, control, accessibility, and failure states.
Measurement	Success measure, guardrail, and review date.
Rollback	How to disable, revert, migrate, or communicate the change.
Documentation	Book section, specification, types, and prompts that must change together.

Change-control rule

A later decision may replace this book, but it must say what it supersedes. New features do not become “approved” merely because they were generated in a chat. A plan accepted in Lovable establishes direction; implementation plus verification establishes current reality; user research establishes evidence; and a marketing statement is permitted only when all three align.

Appendix A. Feature status map

Capability	Status	Evidence / note
Public homepage and discovery CTA	Current	Live published project and index route.
Five-section foundation onboarding	Current	onboarding.tsx and profile schema.
Deterministic pattern recognition	Current	patterns/insights/user-signal logic; no semantic AI.
Tiered career recommendations	Current	recommendations.ts and matching.ts.
Public career explore and filters	Current	explore route and v2.1 filter contract.
Career detail and progression content	Current	Career Detail route and legacy/adapted catalog.
Profile cloud synchronization	Current, limited	Profiles table + local-first sync; JSONB and conflict limitations.
Dashboard journal	Current, isolated	Local entries; does not yet update inference.
Research emotion instrumentation	Current research only	Separate from consumer recommendation logic.
Complete v2.1 career migration	Approved direction	Schema and adapter exist; content migration incomplete.
Approximately 75 balanced careers	Approved direction	Catalog expansion target, contingent on governance.
Discovery Sessions	Approved direction	Progressive Discovery Engine v1.
Evidence ledger and profile versions	Approved direction / prerequisite	Required to make progressive claims true.
Trait-scoped AI conversations	Approved direction	Lovable plan approved; not implemented in supplied code.
Personalized career roadmap	Long-term vision	Current progression content is not a persisted roadmap.
Skill-gap and learning guidance	Long-term vision	Requires stronger skills and resource governance.
Resume/portfolio support	Long-term vision	Requires document privacy and no-fabrication controls.
Mentor matching/community	Hold	Operational trust and safety not defined.
General AI career coach	Hold	Prove contextual conversation pattern first.
B2B/enterprise offering	Open strategy	Consent, neutrality, data separation, and buyer/user tension unresolved.

Appendix B. Source register and authority

ID	Source	Authority in this book
S1	Live Lovable project knowledge and project metadata, accessed 6 Aug 2026.	Current positioning and project identity.
S2	Recent Lovable project messages through 6 Aug 2026.	Recent approved scope and implementation plans; not proof of shipment.
S3	myvocari-main repository snapshot supplied 6 Aug 2026.	Primary evidence for current code, architecture, routes, migrations, and catalog.
S4	Career Compass Project Context & Handoff.	Historical vision, principles, audience, and open questions.
S5	Product Experience Design & Innovation, Group 2, Module 10 final compilation.	Research convergence, limitations, design recommendations, and roadmap evidence.

ID	Source	Authority in this book
S6	Product Experience Design & Innovation, Group 2, Module 7.	Earlier research and prototype evidence incorporated into final compilation.
S7	Progressive Discovery Engine v1.	Approved philosophy and implementation direction for ongoing discovery.
S8	Progressive Discovery Engine v1.1.	Approved philosophy for collaborative profile refinement.
S9	Vocari Career Content & Data Specification v2.1.	Canonical career schema, taxonomy, governance, and recommendation semantics.
S10	Career Content & Data Specifications v1.0 and v2.0.	Historical evolution; superseded by v2.1 where inconsistent.
S11	Module 5 group assignment.	Earlier course framing and conceptual evidence; secondary to final compilation.

Appendix C. Final product test

A future iteration is aligned with Vocari only if the answer to all six questions is yes:

13. Does it improve self-understanding or the quality of an informed next step?
14. Can the user see why Vocari reached the interpretation?
15. Can the user correct, decline, or remove the evidence or change?
16. Is the claim proportionate to the catalog, algorithm, research, and architecture?
17. Does it preserve realistic tradeoffs and avoid deterministic career advice?
18. Will the same concept mean the same thing in copy, data, code, analytics, and AI instructions?

NORTH STAR DECISION

Build the evidence loop before expanding the feature list. Vocari becomes distinctive when a user can contribute a small piece of their story, see a careful interpretation grounded in evidence, correct it when necessary, and receive a more useful next step in return.

End of book • Next formal review: after v2.1 career migration or before the first production AI conversation release, whichever occurs first.